

# Problem Set 1: Solutions

Stat 154, Spring 2018, Prof. Sanchez

*Due date: Tue Feb-06 (before midnight)*

## General Self-Grading Instructions

- Please see the solutions posted on bCourses (in the *Files* folder).
- You will have to enter your score and comments in the *Assingments Comments* section of the correspoind assignment HW01.
- Enter your own scores and comments for (every part) of every problem in the homework on a simple coarse scale:
  - **0** = Didn't attempt or very very wrong,
  - **2** = Got started and made some progress, but went off in the wrong direction or with no clear direction,
  - **5** = Right direction and got half-way there,
  - **8** = Mostly right but a minor thing missing or wrong,
  - **10** = 100% correct.
- Also, enter your total score (out of an overall score of 130 points).
- If there is a cascading error, use a single deduction (don't double or triple or multiple penalize).
- Note: You must justify every self-grade score with a comment. If you are really confused about how to grade a particular problem, you should post on [Piazza](#). This is not supposed to be a stressful process.
- Your self-grades will be due four days after the homework deadline at 11:59 PM sharp (i.e Tue Feb-06).
- We will accept late self-grades up to a week after the original homework deadline for half credit on the associated homework assignment.
- If you don't enter a proper grade by this deadline, you are giving yourself a zero on that assignment.
- Merely doing the homework is not enough, you must do the homework; turn it in on time; read the solutions; do the self-grade; and turn it in on time. Unless all of these steps are done, you will get a zero for that assignment.

## Problem 1 (10 pts)

Create the following matrices in R (and display them).

$$\mathbf{X} = \begin{bmatrix} 2 & 3 \\ -1 & 4 \end{bmatrix}; \quad \mathbf{Y} = \begin{bmatrix} 2 & 0 & 1 \\ 1 & -2 & 3 \end{bmatrix}; \quad \mathbf{Z} = \begin{bmatrix} 1 & 1 \\ -1 & 1 \\ 0 & 2 \end{bmatrix}; \quad \mathbf{W} = \begin{bmatrix} 1 & 0 \\ 8 & 3 \end{bmatrix}; \quad \mathbf{I} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

It is possible to create matrices **X**, **Y**, **Z**, **W**, and **I** in other ways

```
X <- matrix(c(2, -1, 3, 4), nrow = 2)
X
```

```
##      [,1] [,2]
## [1,]    2    3
## [2,]   -1    4
```

```
Y <- matrix(c(2, 1, 0, -2, 1, 3), nrow = 2)
Y
```

```
##      [,1] [,2] [,3]
## [1,]    2    0    1
## [2,]    1   -2    3
```

```
Z <- matrix(c(1, -1, 0, 1, 1, 2), nrow = 3)
Z
```

```
##      [,1] [,2]
## [1,]    1    1
## [2,]   -1    1
## [3,]    0    2
```

```
W <- matrix(c(1, 8, 0, 3), nrow = 2)
W
```

```
##      [,1] [,2]
## [1,]    1    0
## [2,]    8    3
```

```
I <- diag(nrow = 2)
I
```

```
##      [,1] [,2]
## [1,]    1    0
## [2,]    0    1
```

## Problem 2 (10 pts)

Use the matrices created in Problem 1 to perform each of the following operations in R. If the indicated operation cannot be performed, explain why.

a.  $\mathbf{X} + \mathbf{Y}$

Cannot be performed because **X** and **Y** have different number of columns

```
X + Y
```

```
## Error in X + Y: non-conformable arrays
```

b.  $\mathbf{X} + \mathbf{W}$

```
X + W
```

```
##      [,1] [,2]
## [1,]    3    3
## [2,]    7    7
```

c.  $\mathbf{X} - \mathbf{I}$

```
X - I
```

```
##      [,1] [,2]
## [1,]    1    3
## [2,]   -1    3
```

d.  $\mathbf{XY}$

```
X %*% Y
```

```
##      [,1] [,2] [,3]
## [1,]    7   -6   11
## [2,]    2   -8   11
```

e.  $\mathbf{XI}$

```
X %*% I
```

```
##      [,1] [,2]
## [1,]    2    3
## [2,]   -1    4
```

f.  $\mathbf{X} + (\mathbf{Y} + \mathbf{Z})$

Cannot be performed because  $\mathbf{Y}$  and  $\mathbf{Z}$  have different dimensions; likewise  $\mathbf{X}$  and  $\mathbf{Y}$  have different number of columns.

```
X + (Y + Z)
```

```
## Error in Y + Z: non-conformable arrays
```

g.  $\mathbf{Y}(\mathbf{I} + \mathbf{W})$

Cannot be performed because  $\mathbf{Y}$  and  $(\mathbf{I} + \mathbf{W})$  have different number of columns

```
Y %*% (I + W)
```

```
## Error in Y %*% (I + W): non-conformable arguments
```

### Problem 3 (10 pts)

Determine whether the following statements are True or False.

- a. Every orthogonal matrix is nonsingular.

TRUE

- b. Every nonsingular matrix is orthogonal.

FALSE

- c. Every matrix of full rank is square.

FALSE

- d. Every square matrix is of full rank.

FALSE

- e. Every nonsingular matrix is of full rank.

TRUE

### Problem 4 (10 pts)

Let  $\mathbf{X}$ ,  $\mathbf{Y}$ , and  $\mathbf{Z}$  be conformable. Using the properties of transposes, prove that:

$$(\mathbf{XYZ})^T = \mathbf{Z}^T \mathbf{Y}^T \mathbf{X}^T$$

Note that  $(\mathbf{XYZ})^T = [(\mathbf{XY})\mathbf{Z}]^T$

By property of transposes,  $[(\mathbf{XY})\mathbf{Z}]^T = \mathbf{Z}^T (\mathbf{XY})^T$

Likewise,  $(\mathbf{XY})^T = \mathbf{Y}^T \mathbf{X}^T$

Thus:  $(\mathbf{XYZ})^T = \mathbf{Z}^T \mathbf{Y}^T \mathbf{X}^T$

### Problem 5 (10 pts)

Consider the eigenvalue decomposition of a symmetric matrix  $\mathbf{A}$ . Prove that two eigenvectors  $\mathbf{v}_i$  and  $\mathbf{v}_j$  associated with two distinct eigenvalues  $\lambda_i$  and  $\lambda_j$  of  $\mathbf{A}$  are mutually orthogonal; that is,  $\mathbf{v}_i^T \mathbf{v}_j = 0$

By hypothesis, we have:

$$\mathbf{A}\mathbf{v}_i = \lambda_i \mathbf{v}_i \tag{1}$$

and

$$\mathbf{A}\mathbf{v}_j = \lambda_j \mathbf{v}_j \tag{2}$$

with  $\lambda_i \neq \lambda_j$

Taking the transpose of the first equation we have:

$$\mathbf{v}_i^T \mathbf{A} = \lambda_i \mathbf{v}_i^T$$

Premultiplying both sides of equation (2) by  $\mathbf{v}_i^T$  yields

$$\mathbf{v}_i^T \mathbf{A} \mathbf{v}_j = \lambda_j \mathbf{v}_i^T \mathbf{v}_j$$

Substituting  $\mathbf{v}_i^T \mathbf{A}$  by  $\lambda_i \mathbf{v}_i^T$  we have:

$$\lambda_i \mathbf{v}_i^T \mathbf{v}_j = \lambda_j \mathbf{v}_i^T \mathbf{v}_j$$

or

$$(\lambda_i - \lambda_j)(\mathbf{v}_i^T \mathbf{v}_j) = 0$$

Since  $\lambda_i - \lambda_j \neq 0$  by hypothesis, it follows that  $\mathbf{v}_i^T \mathbf{v}_j = 0$

## Problem 6

Refer to the Gram-Schmidt orthonormalization process described in the following wikipedia entry:

[https://en.wikipedia.org/wiki/Gram%E2%80%93Schmidt\\_process](https://en.wikipedia.org/wiki/Gram%E2%80%93Schmidt_process)

This procedure is a method for orthonormalizing a set of vectors in an inner product space. In other words, it allows you to find an orthogonal basis for a set of vectors.

The *projection operator* is given by:

$$proj_{\mathbf{u}}(\mathbf{v}) = \frac{\langle \mathbf{u}, \mathbf{v} \rangle}{\langle \mathbf{u}, \mathbf{u} \rangle} \mathbf{u}$$

This projector operator projects the vector  $\mathbf{v}$  orthogonally onto the line spanned by vector  $\mathbf{u}$ .

### 6.1 Function `inner_product` (10 pts)

Write an R function `inner_product()` that calculates the inner product  $\langle \mathbf{u}, \mathbf{v} \rangle$  of two vectors (of the same length)  $\mathbf{u}$  and  $\mathbf{v}$ .

Given two vectors  $\mathbf{v}$  and  $\mathbf{u}$ , you should be able to invoke your function like:

```
inner_product(v, u)
```

Test `inner_product(v, u)` with  $\mathbf{v} = (1, 3, 5)$  and  $\mathbf{u} = (1, 2, 3)$

```
# function: inner product
# notice that the function checks that 'u' and 'v'
# have equal lengths (this is not mandatory)
inner_product <- function(v, u) {
  if (length(v) != length(u)) {
    stop("\narguments have different lengths")
  }
  vu <- sum(v * u)
  return(vu)
}

v <- c(1, 3, 5)
u <- c(1, 2, 3)

inner_product(v, u)

## [1] 22
```

## 6.2 Function `projection()` (10 pts)

Use your `inner_product()` function to write an R function `projection()` for the projection operator.

Given two vectors  $\mathbf{u}$  and  $\mathbf{v}$ , you should be able to call your function like:

```
projection(v, u)
```

Test `projection(v, u)` with  $\mathbf{v} = (1, 3, 5)$  and  $\mathbf{u} = (1, 2, 3)$

Once again, your function `projection()` may not be exactly as the one below. But the output must match the one in this answer key.

```
# function: projection operator "v onto u"
projection <- function(v, u) {
  uv <- inner_product(u, v)
  uu <- inner_product(u, u)
  proj <- (uv / uu) * u
  return(proj)
}

v <- c(1, 3, 5)
```

```
u <- c(1, 2, 3)
projection(v, u)
```

```
## [1] 1.571429 3.142857 4.714286
```

## Problem 7 (10 pts)

Refer to the same wikipedia entry of the previous question. Once you have your function `projection()`, write R code to apply the Gram-Schmidt orthonormalization procedure to the following sets of vectors:

$\mathbf{x} = (1, 2, 3)$ ;  $\mathbf{y} = (3, 0, 2)$ ;  $\mathbf{z} = (3, 1, 1)$

Start by setting  $\mathbf{u}_1 = \mathbf{x}$ , and report the set of vectors  $\mathbf{u}_k$  and the orthonormalized vectors  $\mathbf{e}_k$ , for  $k = 1, 2, 3$ .

```
# a)
x <- c(3, 2, 1)
y <- c(1, 0, 1)
z <- c(1, 5, 10)

u1 <- x
e1 <- u1 / sqrt(sum(u1 * u1))

u2 <- y - projection(y, u1)
e2 <- u2 / sqrt(sum(u2 * u2))

u3 <- z - projection(z, u1) - projection(z, u2)
e3 <- u3 / sqrt(sum(u3 * u3))

# vectors u
u1; u2; u3
```

```
## [1] 3 2 1
```

```
## [1] 0.1428571 -0.5714286 0.7142857
```

```
## [1] -4.666667 4.666667 4.666667
```

```
# vectors e
e1; e2; e3
```

```
## [1] 0.8017837 0.5345225 0.2672612
```

```
## [1] 0.1543033 -0.6172134 0.7715167
```

```
## [1] -0.5773503 0.5773503 0.5773503
```

## Problem 8 (10 pts)

The length of a vector  $\mathbf{x} = (x_1, x_2, \dots, x_n)$  in the  $n$ -dimensional real vector space  $\mathbb{R}^n$  is usually given by the Euclidean norm:

$$\|\mathbf{x}\| = (x_1^2 + x_2^2 + \dots + x_n^2)^{1/2}$$

In many situations, the Euclidean distance is insufficient for capturing the actual distances in a given space. The class of  $p$ -norms generalizes the notion of *length* of a vector and it is defined by:

$$\|\mathbf{x}\|_p = (|x_1|^p + |x_2|^p + \dots + |x_n|^p)^{\frac{1}{p}}$$

where  $p$  is a real number  $\geq 1$ .

Write a function `lp_norm()` that computes the  $L_p$ -norm of a vector. This function should take two arguments:

- `x` the input vector
- `p` the value for  $p$
- Give `p` a default value of 1
- Allow the user to specify `p = "max"` to compute the  $L_\infty$ -norm

You should be able to call `lp_norm()` like this:

```
lp_norm(x)           # default p = 1
lp_norm(x, p = 2)
lp_norm(x, p = "max") # L-max norm
```

Here's one possible way to implement `lp_norm()`

```
# function to compute Lp-norms
lp_norm <- function(x, p = 1) {
  if (p == "max") {
    return(max(abs(x)))
  } else {
    return((sum(abs(x)^p))^(1/p))
  }
}
```

## Problem 9 (10 pts)

Use your function `lp_norm()` with the following vectors and values for `p`:

- a. `zero <- rep(0, 10)` and `p = 1`



```
# a)
zero <- rep(0, 10)
lp_norm(zero)
```

```
## [1] 0
```

b. ones <- rep(1, 5) and p = 2

```
# b)
ones <- rep(1, 5)
lp_norm(ones, p = 2)
```

```
## [1] 2.236068
```

c. u <- rep(0.4472136, 5) and p = 2

```
# c)
u <- c(0.1825742, 0.3651484, 0.5477226, 0.7302967)
lp_norm(u, p = 2)
```

```
## [1] 1
```

d. u <- 1:500 and p = 100

```
# d)
u <- 1:500
lp_norm(u, p = 100)
```

```
## [1] 508.5663
```

e. u <- 1:500 and p = "max"

```
# e)
u <- 1:500
lp_norm(u, p = "max")
```

```
## [1] 500
```

## Problem 10 (10 pts)

Consider the eigendecomposition of a square matrix  $\mathbf{A}$ .

- a. Prove that the matrix  $b\mathbf{A}$ , where  $b$  is an arbitrary scalar, has  $b\lambda$  as an eigenvalue, with  $\mathbf{v}$  as the associated eigenvector.

Multiplying both members of  $\mathbf{A}\mathbf{v} = \lambda\mathbf{v}$  by  $b$ , we have:

$$(b\mathbf{A})\mathbf{v} = b(\lambda\mathbf{v}) = (b\lambda)\mathbf{v}$$

- b. Prove that the matrix  $\mathbf{A} + c\mathbf{I}$ , where  $c$  is an arbitrary scalar, has  $(\lambda + c)$  as an eigenvalue, with  $\mathbf{v}$  as the associated eigenvector.

The proof is straightforward:

$$\begin{aligned}(\mathbf{A} + c\mathbf{I})\mathbf{v} &= \mathbf{A}\mathbf{v} + c\mathbf{v} \\ &= \lambda\mathbf{v} + c\mathbf{v} \\ &= (\lambda + c)\mathbf{v}\end{aligned}$$

## Problem 11 (20 pts)

For this problem, use the data set `state.x77` that comes in R.

- a. Select the first five columns of `state.x77` and convert them as a matrix; this will be the data matrix  $\mathbf{X}$ . Let  $n$  be the number of rows of  $\mathbf{X}$ , and  $p$  the number of columns of  $\mathbf{X}$

```
# a)
X <- as.matrix(state.x77[,1:5])
n <- nrow(X)
p <- ncol(X)
```

- b. Create a diagonal matrix  $\mathbf{D} = \frac{1}{n}\mathbf{I}$  where  $\mathbf{I}$  is the  $n \times n$  identity matrix. Display the output of `sum(diag(D))`.

```
# b)
D <- diag(1/n, n)
sum(diag(D))
```

```
## [1] 1
```

- c. Compute the vector of column means  $\mathbf{g} = \mathbf{X}^T\mathbf{D}\mathbf{1}$  where  $\mathbf{1}$  is a vector of 1's of length  $n$ . Display (i.e. print)  $\mathbf{g}$ .

```
# c)
ones <- rep(1, n)
g <- t(X) %*% D %*% ones
g
```

```
##           [,1]
## Population 4246.4200
## Income     4435.8000
## Illiteracy  1.1700
## Life Exp    70.8786
## Murder      7.3780
```

- d. Calculate the mean-centered matrix  $\mathbf{X}_c = \mathbf{X} - \mathbf{1}\mathbf{g}^T$ . Display the output of `colMeans(Xc)`.

```
# d)
Xc <- X - ones %*% t(g)
colMeans(Xc)
```

```
##      Population      Income  Illiteracy      Life Exp      Murder
## -3.637979e-14  7.275958e-13 -5.950795e-16 -5.968559e-15  7.283063e-16
```

- e. Compute the (population) variance-covariance matrix  $\mathbf{V} = \mathbf{X}^T\mathbf{D}\mathbf{X} - \mathbf{g}\mathbf{g}^T$ . Display the output of `V`.

```
# e)
V <- t(X) %*% D %*% X - (g %*% t(g))
V
```

```
##      Population      Income  Illiteracy      Life Exp      Murder
## Population 19533050.0836 559805.1840  287.010600 -399.685612 5550.253240
## Income     559805.1840 370021.8400 -160.428000  275.049920 -511.456400
## Illiteracy  287.0106   -160.4280    0.364100   -0.471882    1.550140
## Life Exp    -399.6856    275.0499   -0.471882    1.765980   -3.792091
## Murder      5550.2532   -511.4564    1.550140   -3.792091   13.354916
```

- f. Let  $\mathbf{D}_{1/S}$  be a  $p \times p$  diagonal matrix with elements on the diagonal equal to  $1/S_j$ , where  $S_j$  is the standard deviation for the  $j$ -th variable. Display only the elements in the diagonal of  $\mathbf{D}_{1/S}$

```
# f)
stdevs <- apply(X, 2, function(x) sqrt((n-1)/n) * sd(x))
stdevs
```

```
##      Population      Income  Illiteracy      Life Exp      Murder
## 4419.621034  608.294205    0.603407    1.328902    3.654438
```

```
Ds <- diag(1/stdevs)
```

Ideally, you should obtain the standard deviation dividing by  $\sqrt{n}$ ; if you used `sd()` keep in mind that R calculates it dividing by  $\sqrt{n-1}$  and thus your solutions will slightly vary from the answer key.

- g. Compute the matrix of standardized data  $\mathbf{Z} = \mathbf{X}_c \mathbf{D}_{1/S}$ . Display the output of `colMeans(Z)` and `apply(Z, 2, var)`

```
# g)
Z <- Xc %*% Ds
colMeans(Z)

## [1] -1.755540e-17  1.191824e-15 -9.642287e-16 -4.497301e-15  1.920686e-16

apply(Z, 2, var)

## [1] 1.020408 1.020408 1.020408 1.020408 1.020408
```

Again, recall that R calculates variance dividing by  $n - 1$ .

- h. Compute the (population) correlation matrix  $\mathbf{R} = \mathbf{D}_{1/S} \mathbf{V} \mathbf{D}_{1/S}$ . Display the matrix R

```
# h)
R <- Ds %*% V %*% Ds
R

##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,]  1.00000000  0.2082276  0.1076224 -0.06805195  0.3436428
## [2,]  0.20822756  1.0000000 -0.4370752  0.34025534 -0.2300776
## [3,]  0.10762237 -0.4370752  1.0000000 -0.58847793  0.7029752
## [4,] -0.06805195  0.3402553 -0.5884779  1.00000000 -0.7808458
## [5,]  0.34364275 -0.2300776  0.7029752 -0.78084575  1.0000000
```

- i. Confirm that  $\mathbf{R}$  can also be obtained as  $\mathbf{R} = \mathbf{Z}^T \mathbf{D} \mathbf{Z}$

```
# i)
R2 <- t(Z) %*% D %*% Z
R2

##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,]  1.00000000  0.2082276  0.1076224 -0.06805195  0.3436428
## [2,]  0.20822756  1.0000000 -0.4370752  0.34025534 -0.2300776
## [3,]  0.10762237 -0.4370752  1.0000000 -0.58847793  0.7029752
## [4,] -0.06805195  0.3402553 -0.5884779  1.00000000 -0.7808458
## [5,]  0.34364275 -0.2300776  0.7029752 -0.78084575  1.0000000
```